nativ

company brain · digital colleagues

WHITEPAPER · SECOND EDITION

Company brain

Knowledge from the minds of your employees: minimum viable context (MVC™). The theory and the practice.

July 2026 · gonativ.nl · info@gonativ.nl

Executive summary

Enterprise AI is failing at scale. Despite enormous investments, more than 80% of AI projects deliver no demonstrable business value. Almost all generative AI pilots never reach production.

The technology is not to blame. The problem sits elsewhere: in the minds of your employees. Eighty percent of what an organization knows has never been written down. It lives in the experience of your sales manager, the customer knowledge of your account manager, the exceptions everyone knows but nobody has captured. Today's AI systems — however advanced — can only work with what exists. What is not captured stays invisible.

In the first edition of this whitepaper we introduced the Minimum Viable Context™ (MVC™) framework: a practical method to make the right organizational knowledge usable for AI. That knowledge comes from three sources: existing systems, the minds of employees, and current external data. The biggest blind spot sits in source two — AI cannot work with what was never captured.

This second edition shows how that framework works in practice — as we built it in our platform InsightFlow™, and as we apply it with our clients — a training company and a healthcare organization — alongside active conversations in manufacturing and greentech, among others. The heart of the platform is the company brain: not a pile of documents, but individual, governed facts, organized into eleven knowledge areas. Every fact has an owner, a refresh rhythm, and access rules. And the most important rule of all: AI may propose knowledge, but a human decides what becomes company truth.

On that foundation, digital colleagues do their work — for sales, marketing, and finance, with HR as the fourth field, and as custom colleagues for problems specific to your business. The approach is the same every time: first the knowledge the colleague needs, then the output it must produce, then the tools and the approval moments. A company-specific problem is not a separate adventure — it is the same recipe.

And the brain does not stand still. If you choose to, it grows with the interactions themselves: with explicit consent — fully privacy-compliant — the brain is present at meetings and email traffic, so that every interaction adds to the intelligence of the brain. Everything around a relationship becomes available in one place, and before every conversation the through-line of the past months is ready. The customer decides when the brain reads along — and when it does not.

The science is clear: better results come from better context, not from a better model.

In this whitepaper you will read how to approach that — step by step.

1. Why AI projects fail

1.1 The numbers do not lie

We are investing massively in AI. And yet: most projects deliver nothing. Not because the technology does not work — but because organizations forget what AI actually needs.

NUMBER	WHAT IT MEANS
>80%	of enterprise AI projects deliver no demonstrable value (RAND, 2025)
95%	of generative AI pilots never reach production (MIT Sloan, 2025)
97%	of companies fail to scale AI agents beyond the pilot phase (IDC, 2025)
70%	of AI agent tasks fail in multi-step scenarios (Carnegie Mellon & Salesforce, 2025)

Companies do not fail because they picked the wrong model. They fail because their AI systems lack the knowledge infrastructure to work effectively inside their specific organization.

We know this first-hand. In client conversations we hear the same two sentences almost weekly: "our information is scattered everywhere, nobody knows where anything is" — even at a large technical company where everything seemed meticulously organized. And: "we cannot keep paying for every AI tool" — sales uses one subscription, marketing another, finance a third, with no shared foundation.

1.2 The real problem: knowledge that was never captured

Research points to five root causes of AI failure: a misunderstood problem statement, inadequate data, technology-first thinking, insufficient infrastructure, and problem complexity. The common thread through all of them: a lack of organizational context.

Philosopher Michael Polanyi put it into words back in 1966: "We know more than we can tell." Organizational knowledge exists on a spectrum — from explicit (documented, transferable) to tacit (experience-bound, intuitive, hard to capture).

- 80% of enterprise data is 'dark data' — collected but never analyzed (IBM, 2023)
- 55% of organizational data is unused or unknown to the organization itself (Splunk, 2024)
- Only 14% of enterprise data is active and business-critical (Datamation EMEA, 2024)

The deeper problem is knowledge that was never captured: the decision rules of your sales manager, the customer sensitivities your account manager carries, the exceptions everyone knows but nobody wrote down.

> FROM THE FIELD
>
> At one of our prospects, the employee who does all the custom work is about to leave. There is a neat handover document. But it captures the facts, not the judgment: how she thinks, what she watches for, which pitfalls she avoids. Exactly the knowledge that makes the difference is about to walk out the door.

1.3 The problem with RAG architectures

Retrieval-Augmented Generation (RAG) is the standard approach to grounding AI in organizational knowledge. It retrieves relevant documents and injects them into the context window. Google's ICLR 2025 study confirms: context quality is the most important driver of AI accuracy — weighing heavier than model quality.

But RAG has one critical limitation: it can only retrieve what is already there. The most advanced search engine cannot find knowledge that does not exist.

Retrieval remains a necessary building block — in our own platform too. But the value of retrieval is determined by what there is to find. Most AI infrastructure companies optimize retrieval. Nativ tackles the upstream problem: making sure the right knowledge exists in the first place — and that you can trust it.

> KEY POINT
>
> AI projects do not fail on the model — they fail on **missing organizational context**. The most advanced search engine cannot find knowledge that does not exist.

2. From unwritten knowledge to context engineering

2.1 The tacit knowledge problem

Nonaka and Takeuchi described in 1995 how knowledge moves between tacit and explicit through four steps: socialization (tacit → tacit), externalization (tacit → explicit), combination (explicit → explicit), and internalization (explicit → new tacit). For AI systems, externalization is the bottleneck.

Traditional forms of externalization — documentation, process maps, knowledge bases — are expensive, time-consuming, and suffer from low adoption. Employees resist because it interrupts their workflow and rarely benefits them directly.

2.2 How to actually surface knowledge

The field of knowledge engineering has developed a range of methods: structured interviews, laddering (iterative probing), the critical incident technique, and think-aloud protocols. They work — but they scale poorly. In an organization of 50 employees this can cost hundreds of hours of facilitated sessions with specialized knowledge engineers.

With the arrival of AI, that has changed.

Nativ broke this bottleneck with InsightFlow™. A digital interviewer conducts structured video interviews with employees, built along the knowledge areas. Every interview is processed into a reviewable report — a distilled answer per question — that is only added to the brain after human approval. For larger groups we deploy studies: structured research across large groups of participants at once, with questions drafted from what the company already knows — so you never ask questions the company can already answer.

What used to cost hundreds of hours of specialized knowledge engineers now happens in a structured way, at scale.

In short: the knowledge AI needs already exists — but largely in people's heads. It is not a model problem. It is a problem of uncaptured knowledge. The question is how to extract that knowledge from those heads in a structured way, capture it, and offer it to an AI system at the right moment. That is context engineering.

2.3 Context engineering as a discipline

In 2025-2026, context engineering grew into a recognized discipline. Tobi Lütke (Shopify) called it 'the #1 skill of 2026.' Where prompts are short task descriptions for everyday use, Andrej Karpathy describes the industrial discipline:

> *"Context engineering is the delicate art and science of filling the context window with just the right information for the next step."*
>
> — Andrej Karpathy, 2025

2.4 Context beats model quality

Multiple studies confirm it: context quality outweighs model quality for enterprise AI performance.

- Google (ICLR 2025): a smaller model with better context beats a larger model with worse context.
- Glean reports 50% productivity improvement and 250–422% ROI with context-engineered agents.
- California Management Review (2026): organizations that structure tacit knowledge see an 80% reduction in expert workload.

The highest-leverage investment for enterprise AI is not a better model. It is better context.

3. The Nativ MVC Framework™

3.1 What is Minimum Viable Context™?

Just as a Minimum Viable Product (MVP) defines the smallest product that delivers value, Nativ's Minimum Viable Context™ (MVC™) defines the minimal organizational knowledge an AI agent needs to perform a task reliably.

The framework builds on decades of knowledge management research — from Polanyi and Nonaka to current analyses by RAND, IDC, MIT Sloan, and Gartner. It follows the Pareto principle applied to organizational knowledge: a small, deliberate subset of context delivers the majority of AI agent capability. In our first implementations we see this pattern: with roughly twenty percent of the knowledge you cover eighty percent of the tasks. Not everything at once — but the right things first.

That is also the answer to the most common objection we hear in client conversations: "our data is not ready for this." You do not need a data migration project first. You start with the right twenty percent.

MVC™ is the method. The company brain in InsightFlow™ is the execution. The rest of this whitepaper describes that execution as it runs — not as it was imagined on a drawing board.

3.2 Three sources, one brain

A company brain is built from three kinds of sources, depending on where the knowledge lives.

Existing systems. When knowledge is already stored somewhere, we connect it through ready-made integrations: CRM, email, meeting notes, accounting, advertising and social media platforms, the company website, and uploaded documents. Nativ does this too — provided the data actually exists and is in good shape.

The minds of employees. When knowledge is not captured, we extract it directly from employees' heads: through AI interviews with a digital interviewer and through studies across larger groups. This is the part no integration can solve. You cannot search what was never written down.

External sources. For market and competitive information we deploy research agents that draw on around ten live data sources — from web search engines to news and social media data. At the moment it matters, or on a fixed rhythm. Not as a bulk import, but as a signal.

The art is not in choosing one source. It is in combining them intelligently per use case.

3.3 Facts, not files

The biggest difference with a classic knowledge base sits in the smallest unit. The company brain does not store knowledge as documents, but as individual facts: "our payment term is 30 days", "client X does not want phone contact", "we no longer sell this service".

That sounds like a detail. It is the foundation. A document can be half right; a fact is either right or wrong. Because every fact is its own unit, the brain can track per fact who is responsible for it, how current it is, and who may see it. Every fact is also automatically made available in both Dutch and English — your assistant switches effortlessly between the two.

3.4 Full context: a brain that grows with you

Interviews and studies capture what lives in people's heads today. But an organization keeps talking every day — in meetings, in email, with customers. That is why the company brain is designed for **full context**: if the customer chooses to, the brain is — with explicit consent, fully privacy-compliant — present at those interactions, so that every interaction adds to the intelligence of the brain.

You notice the difference in practice. All interactions around a relationship are bundled — emails and meetings linked to the person they are about. Anyone walking into a meeting can pull up, at any moment, the through-line of the conversations and email exchanges of, say, the past two months — and always arrives at the table optimally prepared. For sales this is immediately gold; the same principle works in every field, finance included — payment agreements and commitments live in conversations and email threads too.

This builds something no one-off capture can offer: a growing store of tested knowledge — not what was said once, but what is confirmed, corrected, and sharpened interaction after interaction.

And the customer draws the line. You decide when the brain reads along and when it does not — per team, per situation. It is the same trust boundary as in chapter 5: the brain works for your people, never behind their backs.

> KEY POINT
>
> MVC™ is the method, the company brain the execution: **individual, governed facts** instead of a pile of documents — filled from systems, minds, and external sources, and, if you choose, **growing with every interaction**. The customer decides when the brain reads along.

4. The eleven knowledge areas

In the first edition we described ten context blocks. In practice — after implementations with clients — they became eleven, developed into a template of over four hundred knowledge categories. That template is the starting point of every company brain: it describes which facts an organization should capture, so you never start from a blank page.

11 knowledge areas together form the company brain	400+ knowledge categories in the template — never a blank page	3 kinds of sources: systems, minds, external data

KNOWLEDGE AREA	WHAT IT CAPTURES
1. Identity	Mission, vision, values, story, USPs
2. Products & services	Offering, features, pricing, upsell paths
3. People & roles	Who does what, decision authority, expertise
4. Market & competition	Competitors, trends, your position
5. Brand & tone	Visual identity, writing style, templates, examples
6. Customers	Segments, contacts, contracts, sensitivities
7. Processes	Sales, onboarding, complaint handling, invoicing
8. Finance	Pricing agreements, margins, payment terms, budget rules
9. Tools & systems	Which systems you use, for what, and who manages them
10. Pipeline & projects	Running deals and projects, status, blockers
11. Strategy	Goals, choices, roadmap, key decisions

Unlike the first edition, we assign ownership and refresh rhythm not per area, but per fact. That is more precise: the price list changes quarterly, but the discount agreement with your largest customer can change next week. Every fact gets its own rhythm — more on that in chapter 5.

Every fact is small enough to capture quickly, and the template is complete enough to deliver value immediately. That is the core of MVC™: not everything at once, but the right things first.

5. One brain, governed knowledge

Capturing knowledge is necessary but not sufficient. A brain nobody dares to trust does not get used. That is why every fact in the company brain is governed: it carries its own rules — who owns it, how long it stays valid, and who may see it. Governance does not sit next to the brain; it sits inside it.

5.1 Every fact has an owner and an expiry date

Every fact gets one owner: the person who stands behind it. And every fact has a refresh rhythm — from weekly to never (for facts that do not change, like a founding year). The brain therefore knows at any moment which knowledge is fresh, which needs attention soon, and which has expired.

This is not bookkeeping for its own sake. Expired facts are automatically kept out of AI answers: better an assistant that says "I am not sure about that" than one that confidently quotes an old price. And the health of the brain is visible daily on a dashboard: what percentage is fresh, what has expired, what has no owner. Knowledge quality you can inspect — not assume.

5.2 AI proposes, a human decides

The most important design choice in InsightFlow™ is a trust boundary: AI may propose knowledge, but never declare it company truth on its own.

Everything the platform derives as a business fact — from an interview, a study, or an individual contribution — is staged as a proposal. Only when an employee approves it — traceably, by name — does it become part of the brain. No silent changes, no AI filling its own knowledge base with assumptions.

That keeps the company brain what it must be: the captured truth of your organization — not a model's best guess.

5.3 Not all knowledge is for everyone

An AI system may only answer based on knowledge the asking user would have been allowed to see themselves. That is why every fact carries its own access rules: visible to the whole organization, to specific departments, or to leadership only. Sensitive facts can additionally be marked secret — those are skipped for every AI at retrieval. And facts marked as personal data are kept out of AI communication that leaves the organization.

Enforcement happens at the moment of retrieval: a fact you may not see is not offered to the model — and so cannot be cited as a company fact in an answer. Access rules are applied by the platform itself — not by the user or their device.

The same boundaries apply to what enters the brain in the first place. At a healthcare organization we work with, a hard data boundary applies: medical and patient data never enter the platform — the digital colleagues there will work exclusively on the business layer, from financial figures to procedures and agreements.

This is not optional. According to Gartner, 40% of Copilot rollout projects are delayed by oversharing risks. So we design access in, not around.

5.4 From goal to data source: the four-layer method

To determine which knowledge gets priority, we use a goal-driven mapping in four steps in every implementation:

1. **Goal** — which business result must this agent support?
2. **Tasks** — which specific tasks belong to it?
3. **Decisions** — which information is needed for those?
4. **Data sources** — where does that information live today, and in what form?

This ensures context capture is driven by real business needs — not by abstract taxonomies. Remember this mapping: in chapter 6 you will see that every digital colleague — standard or custom — starts from it.

> KEY POINT
>
> **AI proposes, a human decides.** Every fact carries its own owner, expiry date, and access rules — and expired or restricted knowledge never reaches the model.

6. From company brain to digital colleagues

6.1 The assistant for everyone

The first thing that runs on the company brain is a personal AI assistant for every employee — ours is called Alice. She writes emails in the right tone, answers internal questions, formats documents to brand guidelines, and gets new employees up to speed quickly — in Dutch or English.

A standout feature is the No-Hallucination mode: a setting in which the assistant answers exclusively from approved company facts. If the brain does not know something, the assistant says so — instead of inventing something plausible. And if the brain yields nothing useful, the answer honestly notes that it leans on general knowledge rather than company facts. For a sales question about your own pricing, you want company truth, not probability.

> FROM THE FIELD
>
> At one of our clients, management debated the website for ninety minutes without resolution. When someone put the same question to the digital assistant, a proposal came back immediately that everyone could get behind — fed by the captured brand principles and earlier decisions.

6.2 Digital colleagues per field

For recurring departmental work we go a step further than an assistant that gives answers: digital colleagues that execute work. They cover four fields:

FIELD	COLLEAGUE	WHAT THEY DO
Sales	Hunter	Lead research and client dossiers, meeting preparation, proposals, follow-up
Marketing	Brandon	Campaign planning, content and social posts on brand, image production, performance analysis
Finance	Billy	Receivables and cash monitoring, month-end close, monthly reporting, budget round, price monitoring
HR	Crystal	Onboarding of new employees, leave and HR questions, job postings

Every digital colleague works on the same governed brain — there are no separate knowledge bases per colleague. Each colleague leans mainly on the knowledge areas of their field: sales on customers and pipeline, finance on finance and processes, marketing on brand and market. What a colleague gets to see is determined by the access rules of the employee requesting the work — and for scheduled tasks by the department the task was set up for. And colleagues work together: a colleague preparing a client meeting first automatically has a fresh client dossier built.

> FROM THE FIELD
>
> At a training company (twenty employees, sixty freelance trainers) the company brain runs with digital marketing colleagues across two brands: live within four weeks, a €3,000–4,000-per-month marketing role that had gone unfilled for months now filled for €995 per month, and double the output.
>
> The most important lesson from our implementations is not technical but about adoption. In one marketing team we supported, the data-driven team members embraced the tool immediately; the creatives never opened it on their own. The solution was to reverse the direction of work — the digital colleague prepares the weekly report and the team responds to it. After a few weeks we then keep hearing the same question: "how did I even do this before?"

6.3 Colleagues that account for their work

A digital colleague that communicates externally without control is a risk — not an asset. That is why accountability is built into how colleagues work:

Approval before action. Work that matters waits for human sign-off. Whoever rejects at such an approval moment explains why — and the colleague redoes that step with the feedback, like a junior processing the red marks in a first draft. The monthly financial report even requires two signatures in sequence: first the controller, then the CFO.

Honesty over completeness. Finance colleagues take figures exactly as they appear in the accounting system and supplied exports, and never do their own arithmetic. Every price in a proposal is traceable to its source. And a follow-up colleague that finds no good reason to email says so: better no message than an invented one.

Learning from every correction. Approvals, edits, and rejections are captured and processed into lessons per company, per colleague. The colleague carries those lessons into every next piece of work — adapting, like a new employee, to the way you work. And every piece of work is additionally graded on quality, including by an independent AI review step deliberately decoupled from the colleague that produced the work — so no one is grading their own homework.

6.4 Custom colleagues: the same recipe for your specific problem

Not every problem fits a field. Clients regularly ask us for a colleague for one specific task: summarizing an intake that now costs hours of manual work into a briefing for the person delivering the service, or watching a threshold in the schedule that automatically triggers extra marketing the moment it is reached.

For such custom colleagues the approach is exactly the same as for the standard teams. We start with the goal-driven mapping from chapter 5 and then pin down four things:

1. **What knowledge does this colleague need?** Which facts from the brain, which data sources, which systems.
2. **What must come out?** The concrete output: a briefing, a report, a draft email — in what format, for whom.

3. **Which tools does it use?** Every connection and every action the colleague may take is named and bounded up front.
4. **Where is the human?** The approval moments: what the colleague may do independently, and what waits for sign-off.

A custom colleague then runs on the same brain, under the same access rules, with the same approval moments and the same learning loop as every standard colleague. Not a consultancy adventure, but a repeatable recipe. That means you can predict exactly what you will get — and we solve a company-specific problem just as transparently as a standard task.

> KEY POINT
>
> Digital colleagues across four fields — and custom ones — work on one governed brain, with **approval before action**, traceable figures, and lessons from every correction. Always the same recipe.

7. How we build it

7.1 Phase 1 — Scan

Before we build, we map. The AI Opportunity Scan (duration: one week) identifies where AI makes the most impact and which knowledge gaps exist. We map the organization against the eleven knowledge areas: what is already partly documented? What lives purely in people's heads?

7.2 Phase 2 — The foundation

Every implementation follows a worked-out delivery plan — from signed agreement to working product — with, for every step, whose turn it is: the client or Nativ. Fixed decision moments keep the process tight: light sign-offs where possible, a formally digitally signed agreement where required. And progress is visible to the client at every moment: which step is running, whose turn it is, what has been completed.

The company brain starts from the template in chapter 4: every relevant knowledge area is turned into concrete, open questions. Those are then filled through three channels — integrations with existing systems, AI interviews with employees, and documents that already exist. Every fact derived from them lands as a proposal and passes the human approval from chapter 5.

There are two variants: **configure** (setting up the standard platform — the shortest route, in clear steps) and **build together** (developing custom colleagues — with weekly alignment and roughly four weeks of build time). From go-live we measure the baseline: time saved on routine questions, document creation, and onboarding time.

We deliberately start small and reversible: one department or location first, expanding once usage proves it. And the brain remains yours — whoever stops receives a complete export of their company facts.

FROM THE FIELD

At a healthcare organization with multiple locations, the rollout deliberately starts at one location, with a short commitment and a data boundary agreed up front. Only when usage proves itself there does the rest of the organization follow. The decision stays with the client — at every step.

7.3 Phase 3 — Digital colleagues per department

Based on measured value we add digital colleagues step by step — starting with the department with the highest expected impact (usually sales or marketing). Every colleague starts on the existing brain; per department we enrich the knowledge areas that work requires. We measure the department-specific KPIs, iterate, and expand.

7.4 Continuous maintenance

Every fact has its own refresh rhythm, and the brain enforces it on its own: owners see in their personal work queue exactly which of their facts need attention, and the daily health dashboard shows how fresh the brain is. The result is a sustainable maintenance system instead of a one-off project. The company brain improves continuously as knowledge accumulates and is refined.

And because we use our own platform daily — our own company brain is colleague number four — we experience every rough edge before our clients do.

8. Why this is different

8.1 Beyond document search

The dominant approach to enterprise AI knowledge management is document-centric: index existing documents, emails, and databases, then retrieve relevant passages. Tools like Glean, Notion AI, and enterprise search platforms do this. Valuable — but limited to knowledge that is already written down somewhere.

Nativ addresses the 80% that is not written down. The company brain does not just search existing knowledge — it creates new organizational knowledge by converting tacit expertise into structured, governed, retrievable facts.

8.2 Beyond integration platforms

MCP servers and comparable integration tools connect AI agents to existing software and data sources. That solves the access problem — but not the existence problem. If knowledge lives in someone's head and sits in no system, no integration can surface it.

Moreover, a connection does not equal trust. Our integrations feed a governed brain — with ownership, freshness, and human approval — and an execution layer in which digital colleagues work under supervision. A pass-through is not a colleague.

8.3 The Nativ position

Nativ occupies a unique position in the enterprise AI stack: knowledge creation infrastructure. That position rests on three verbs:

- **Create** — surfacing knowledge that was never written down anywhere: interviews, studies, research — and, with explicit consent, the daily interactions in which knowledge is born.
- **Govern** — every fact with an owner, an expiry date, and access rules, and a human who decides what becomes truth.
- **Execute** — digital colleagues that do real work on that knowledge, with approval where it counts — from standard team to custom colleague, always the same recipe.

Where others optimize the retrieval of existing knowledge, Nativ generates and governs the knowledge that makes retrieval valuable — and puts digital colleagues to work on it.

This is the missing infrastructure layer that explains why so many RAG implementations underperform: they search a knowledge base that contains only a fraction of the organization's real intelligence.

9. Conclusion

The enterprise AI failure crisis is not a technology problem. It is a knowledge infrastructure problem.

Context engineering — the systematic discipline of identifying, capturing, structuring, and delivering organizational knowledge to AI systems — is the missing layer.

The Nativ MVC Framework™ offers a practical, scientifically grounded method to build a company brain from three kinds of knowledge sources: existing systems, the minds of employees, and external sources. In InsightFlow™ that method became a product: eleven knowledge areas with over four hundred categories, every fact with an owner, a refresh rhythm, and access rules — and a human who decides what becomes company truth. On top of it work digital colleagues for sales, marketing, and finance — with HR as the fourth field — and custom colleagues for the problems specific to your business, built with the same recipe. And for those who choose it, the brain grows with the interactions themselves — with explicit consent and under your own control — so that knowledge is not captured once, but continuously confirmed and enriched. It runs in practice today: live at a training company, rolling out at a healthcare organization.

The evidence is clear: context quality outweighs model quality. The organizations that truly win with AI are not the ones with the biggest models. They are the organizations that capture and govern what their people know.

Because in the era of AI agents, the uncaptured knowledge of your organization is your last defensible competitive advantage.

References

Cleanlab (2025). AI Leaders Survey. · Datamation (2024). EMEA Data Classification Study. · Ericsson & Simon (1993). Protocol Analysis. MIT Press. · Flanagan (1954). Critical incident technique. Psych. Bulletin, 51(4). · Gartner (2026). Context Engineering: Definition. · Geiwitz et al. (1990). Expert System for KA Techniques. · Glean (2025). Context Engineering Report. · Google Research (2025). Retrieval-augmented generation. ICLR 2025. · Hoffman (1987). Extracting expert knowledge. AI Magazine, 8(2). · IBM (2023). Dark Data. · IDC (2025). AI Agent Scalability Survey. · Karpathy (2025). Context engineering in LLM applications. · MIT Sloan (2025). Why 95% of GenAI Pilots Fail. · Nonaka & Takeuchi (1995). The Knowledge-Creating Company. · Polanyi (1966). The Tacit Dimension. · RAND (2025).